

A Producer Library Interface to DWARF

David Anderson

1. INTRODUCTION

This document describes an interface to `libdwarf`, a library of functions to provide creation of DWARF debugging information records, DWARF line number information, DWARF address range and pubnames information, weak names information, and DWARF frame description information.

1.1 Copyright

Copyright 1993-2006 Silicon Graphics, Inc.

Copyright 2007-2021 David Anderson.

Permission is hereby granted to copy or republish or use any or all of this document without restriction except that when publishing more than a small amount of the document please acknowledge Silicon Graphics, Inc and David Anderson.

This document is distributed in the hope that it would be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

1.2 Purpose and Scope

The purpose of this document is to propose a library of functions to create DWARF debugging information. Reading (consuming) of such records is discussed in a separate document.

The functions in this document have mostly been implemented at Silicon Graphics and used by the SGI code generator to provide DWARF2 debugging information in the 1990's. Some functions (and support for some extensions) were provided by Sun Microsystems.

Example code showing one use of the functionality may be found in the `dwarfgen` `dwarfgen` and `simpleexample` application (provided in the source distribution along with `libdwarf`).

The focus of this document is the functional interface, and as such, implementation and optimization issues are intentionally ignored.

Error handling, error codes, and certain `Libdwarf` codes are discussed in the "*A Consumer Library Interface to DWARF*", which should be read before reading this document.

Before December 2018 very few functions in the Producer Library follow the error-returns as defined in "*A Consumer Library Interface to DWARF*".

As of December 2018 every Producer Library call has a version that supports that Consumer Library Interface and returns `DW_DLV_OK` or `DW_DLV_ERROR` (the Producer Library has no use of `DW_DLV_NO_ENTRY`). The table of contents of this document lists the latest version of each function. However, all the earlier documentation is present here immediately following the documentation of the latest, and preferred, interface. All the earlier interfaces are supported in the library.

Early interfaces (before December 2018) The general style of functions here in the producer library is rather C-traditional with various types as return values (quite different from the consumer library interfaces). The style generally follows the style of the original DWARF1 reader proposed as an interface to DWARF. When the style of the reader interfaces was changed (1994) in the dwarf reader (See the "Document History" section of "*A Consumer Library Interface to DWARF*") the interfaces here were not changed as it seemed like too much of a change for the two applications then using the interface! So this interface remains in the traditional C style of returning various data types with various (somewhat inconsistent) means of indicating failure.

December 2018 and later function interfaces all return either `DW_DLV_OK` or `DW_DLV_ERROR` in a simple int.

The error handling code in the library may either return a value or abort. The library user can provide a function that the producer code will call on errors (which would allow callers avoid testing for error returns if the user function exits or aborts). See the `dwarf_producer_init()` description below for more details.

1.3 Document History

This document originally prominently referenced "UNIX International Programming Languages Special Interest Group " (PLSIG). Both UNIX International and the affiliated Programming Languages Special Interest Group are defunct (UNIX is a registered trademark of UNIX System Laboratories, Inc. in the United States and other countries). Nothing except the general interface style is actually related to anything shown to the PLSIG (this document was open sourced with `libdwarf` in the mid 1990's).

See "<http://www.dwarfstd.org>" for information on current DWARF standards and committee activities.

1.4 Definitions

DWARF debugging information entries (DIEs) are the segments of information placed in the `.debug_info` and related sections by compilers, assemblers, and linkage editors that, in conjunction with line number entries, are necessary for symbolic source-level

debugging. Refer to the document "*DWARF Debugging Information Format*" from UI PLSIG for a more complete description of these entries.

This document adopts all the terms and definitions in "*DWARF Debugging Information Format*" version 2. and the "*A Consumer Library Interface to DWARF*".

In addition, this document refers to Elf, the ATT/USL System V Release 4 object format. This is because the library was first developed for that object format. Hopefully the functions defined here can easily be applied to other object formats.

1.5 Overview

The remaining sections of this document describe a proposed producer (compiler or assembler) interface to *Libdwarf*, first by describing the purpose of additional types defined by the interface, followed by descriptions of the available operations. This document assumes you are thoroughly familiar with the information contained in the *DWARF Debugging Information Format* document, and "*A Consumer Library Interface to DWARF*".

The interface necessarily knows a little bit about the object format (which is assumed to be Elf). We make an attempt to make this knowledge as limited as possible. For example, *Libdwarf* does not do the writing of object data to the disk. The producer program does that.

1.6 Revision History

March 1993	Work on dwarf2 sgi producer draft begins
March 1999	Adding a function to allow any number of trips through the dwarf_get_section_bytes_a() call.
April 10 1999	Added support for assembler text output of dwarf (as when the output must pass through an assembler). Revamped internals for better performance and simpler provision for differences in ABI.
Sep 1, 1999	Added support for little- and cross- endian debug info creation.
May 7 2007	This library interface now cleans up, deallocating all memory it uses (the application simply calls dwarf_producer_finish(dbg)).
September 20 2010	Now documents the marker feature of DIE creation.
May 01 2014	The dwarf_producer_init() code has a new interface and DWARF is configured at run time by its arguments. The producer code used to be configured at configure time, but the configure time producer configure options are no longer used. The configuration was unnecessarily complicated: the run-time configuration is simpler to understand.

- September 10, 2016 Beginning the process of creating new interfaces so that checking for error is consistent across all calls (as is done in the consumer library). The old interfaces are kept and supported so we have binary and source compatibility with old code.
- December 01, 2018 All function interfaces now have a version that returns only DW_DLV_OK or DW_DLV_ERROR and pointer and other values are returned through pointer arguments. For example, dwarf_add_frame_info_c() is the December 2018 version, while dwarf_add_frame_info(), dwarf_add_frame_info_b() are earlier versions.
- July 14, 2020 To enable testing of reading the DWARF5 section .debug_sup the new function dwarf_add_debug_sup() is added. dwarfgen can call this function, though dwarfgen presently only fills out a bogus .debug_sup section to enable simple testing.
- January 25, 2021 dwarf_add_AT_block_a() now also supports the DWARF5 form DW_FORM_exprloc.
- June 2021 Removing the obsolete functions that return Dwarf_Unsigned etc and required ugly casting to check success/fail. The ones returning int DW_DLV_OK etc are the only ones that should be used. The library is now in its own file (libdwarfp.a or libdwarfp.so) and the source in its own directory (src/lib/libdwarfp). libdwarfp is only built if a build of dwarfgen is requested. Meaning with standard builds this library is not provided. It only creates DWARF2 with any completeness, so it's not clear how it could possibly be generally useful.

2. Type Definitions

2.1 General Description

The *libdwarf.h* header file contains typedefs and preprocessor definitions of types and symbolic names used to reference objects of and types used by *libdwarf* and some declarations needed by *libdwarfp*. The *libdwarfp.h* header file defines producer functions and type specifically used by *libdwarfp*. The types defined by typedefs contained in *libdwarf.h* and *libdwarfp.h* all use the convention of adding *Dwarf_* as a prefix to indicate that they refer to objects used by Libdwarf. The prefix *Dwarf_P_* is used for objects referenced by the *Libdwarf* Producer when there are similar but distinct objects used by the Consumer.

2.2 Namespace issues

Application programs should avoid creating names beginning with `Dwarf_ dwarf_` or `DW_` as these are reserved to dwarf and libdwarf.

3. libdwarf and Elf and relocations

Much of the description below presumes that Elf is the object format in use. The library is probably usable with other object formats that allow arbitrary sections to be created. The library does not write anything to disk. Instead it provides access so that callers can do that in whatever object format is appropriate.

3.1 binary or assembler output

With `DW_DLC_STREAM_RELOCATIONS` (see below) it is assumed that the calling app will simply write the streams and relocations directly into an Elf file, without going through an assembler.

With `DW_DLC_SYMBOLIC_RELOCATIONS` the calling app must either A) generate binary relocation streams and write the generated debug information streams and the relocation streams direct to an elf file or B) generate assembler output text for an assembler to read and produce an object file.

With case B) the libdwarf-calling application must use the relocation information to change points of each binary stream into references to symbolic names. It is necessary for the assembler to be willing to accept and generate relocations for references from arbitrary byte boundaries. For example:

```
.data 0a0bcc  #producing 3 bytes of data.  
.word mylabel  #producing a reference  
.word endlabel - startlabel #producing absolute length
```

3.2 libdwarf relationship to Elf

When the documentation below refers to 'an elf section number' it is really only dependent on getting (via the callback function passed by the caller of `dwarf_producer_init()`. a sequence of integers back (with 1 as the lowest).

When the documentation below refers to 'an Elf symbol index' it is really dependent on Elf symbol numbers only if `DW_DLC_STREAM_RELOCATIONS` are being generated (see below). With `DW_DLC_STREAM_RELOCATIONS` the library is generating Elf

relocations and the section numbers in binary form so the section numbers and symbol indices must really be Elf (or elf-like) numbers.

With `DW_DLC_SYMBOLIC_RELOCATIONS` the values passed as symbol indexes can be any integer set or even pointer set. All that libdwarf assumes is that where values are unique they get unique values. Libdwarf does not generate any kind of symbol table from the numbers and does not check their uniqueness or lack thereof.

3.3 libdwarf and relocations

With `DW_DLC_SYMBOLIC_RELOCATIONS` libdwarf creates binary streams of debug information and arrays of relocation information describing the necessary relocation. The Elf section numbers and symbol numbers appear nowhere in the binary streams. Such appear only in the relocation information and the passed-back information from calls requesting the relocation information. As a consequence, the 'symbol indices' can be any pointer or integer value as the caller must arrange that the output deal with relocations.

With `DW_DLC_STREAM_RELOCATIONS` all the relocations are directly created by libdwarf as binary streams (libdwarf only creates the streams in memory, it does not write them to disk).

3.4 symbols, addresses, and offsets

The following applies to calls that pass in symbol indices, addresses, and offsets, such as `dwarf_add_AT_targ_address_c()` `dwarf_add_arange_c()` and `dwarf_add_frame_fde_c()`.

With `DW_DLC_STREAM_RELOCATIONS` a passed in address is one of: a) a section offset and the (non-global) symbol index of a section symbol. b) A symbol index (global symbol) and a zero offset.

With `DW_DLC_SYMBOLIC_RELOCATIONS` the same approach can be used, or, instead, a passed in address may be c) a symbol handle and an offset. In this case, since it is up to the calling app to generate binary relocations (if appropriate) or to turn the binary stream into a text stream (for input to an assembler, if appropriate) the application has complete control of the interpretation of the symbol handles.

4. Memory Management

Several of the functions that comprise the *Libdwarf* producer interface dynamically allocate values and some return pointers to those spaces. The dynamically allocated spaces can not be reclaimed (and must not be freed) except that all such libdwarf-allocated memory is freed by `dwarf_producer_finish_a(dbg)` .

All data for a particular Dwarf_P_Debug descriptor is separate from the data for any other Dwarf_P_Debug descriptor in use in the library-calling application.

4.1 Read Only Properties

The read-only properties specified in the consumer interface document do not generally apply to the functions described here.

4.2 Storage Deallocation

Calling `dwarf_producer_finish_a(dbg)` frees all the space, and invalidates all pointers returned from `Libdwarf` functions on or descended from `dbg`).

4.3 Error Handling

In general any error detected by the producer should be considered fatal. That is, it is impossible to produce correct output so producing anything seems questionable.

The original producer interfaces tended to return a pointer or a large integer as a result and required the caller to cast that value to determine if it was actually a -1 meaning there was an error.

Beginning in September 2016 additional interfaces are being added to eliminate the necessity for callers to do this ugly casting of results. In December 2018 that process has reached completion. The revised functions return `DW_DLV_OK`, or `DW_DLV_ERROR`. (which are small signed integers) and will have an additional pointer argument that will provide the value that used to be the return value. This will make the interfaces type-safe.

The function `dwarf_get_section_bytes_a()` can also return `DW_DLV_NO_ENTRY`.

5. Functional Interface

This section describes the functions available in the *Libdwarf* library. Each function description includes its definition, followed by a paragraph describing the function's operation.

The following sections describe these functions.

The functions may be categorized into groups: *initialization and termination operations*, *debugging information entry creation*, *Elf section callback function*, *attribute creation*, *expression creation*, *line number creation*, *fast-access (aranges) creation*, *fast-access (pubnames) creation*, *fast-access (weak names) creation*, *macro information creation*, *low level (.debug_frame) creation*, and *location list (.debug_loc) creation*.

5.1 Initialization and Termination Operations

These functions setup `Libdwarf` to accumulate debugging information for an object,

usually a compilation-unit, provided by the producer. The actual addition of information is done by functions in the other sections of this document. Once all the information has been added, functions from this section are used to transform the information to appropriate byte streams, and help to write out the byte streams to disk.

Typically then, a producer application would create a `Dwarf_P_Debug` descriptor to gather debugging information for a particular compilation-unit using `dwarf_producer_init()`.

The producer application would use this `Dwarf_P_Debug` descriptor to accumulate debugging information for this object using functions from other sections of this document. Once all the information had been added, it would call `dwarf_transform_to_disk_form_a()` to convert the accumulated information into byte streams in accordance with the DWARF standard. The application would then repeatedly call `dwarf_get_section_bytes_a()` for each of the `.debug_*` created. This gives the producer information about the data bytes to be written to disk. At this point, the producer would release all resource used by `Libdwarf` for this object by calling `dwarf_producer_finish_a()`.

It is also possible to create assembler-input character streams from the byte streams created by this library. This feature requires slightly different interfaces than direct binary output. The details are mentioned in the text.

5.1.1 `dwarf_producer_init()`

```
int dwarf_producer_init(  
    Dwarf_Unsigned flags,  
    Dwarf_Callback_Func func,  
    Dwarf_Handler errhand,  
    Dwarf_Ptr errarg,  
    void * user_data  
    const char *isa_name,  
    const char *dwarf_version,  
    const char *extra,  
    Dwarf_P_Debug *dbg_returned,  
    Dwarf_Error *error)
```

The function `dwarf_producer_init()` returns a new `Dwarf_P_Debug` descriptor that can be used to add Dwarf information to the object. On success it returns `DW_DLV_OK`. On error it returns `DW_DLV_ERROR`. `flags` determine whether the target object is 64-bit or 32-bit. `func` is a pointer to a function called-back from `Libdwarf` whenever `Libdwarf` needs to create a new object section (as it will for each `.debug_*` section and related relocation section).

The flags values (to be OR'd together in the flags field in the calling code) are as follows:

DW_DLC_WRITE is required. The values DW_DLC_READ DW_DLC_RDWR are not supported by the producer and must not be passed.

The flag bit DW_DLC_POINTER64 (or DW_DLC_SIZE_64) Indicates the target has a 64 bit (8 byte) address size. The flag bit DW_DLC_POINTER32 (or DW_DLC_SIZE_32) Indicates the target has a 32 bit (4 byte) address size. If none of these pointer sizes is passed in DW_DLC_POINTER32 is assumed.

The flag bit DW_DLC_OFFSET32 indicates that 32bit offsets should be used in the generated DWARF. The flag bit DW_DLC_OFFSET64 DW_DLC_OFFSET_SIZE_64 indicates that 64bit offsets should be used in the generated DWARF.

The flag bit DW_DLC_IRIX_OFFSET64 indicates that the generated DWARF should use the early (pre DWARF3) IRIX method of generating 64 bit offsets. In this case DW_DLC_POINTER64 should also be passed in, and the isa_name passed in (see below) should be "irix".

If DW_DLC_TARGET_BIGENDIAN or DW_DLC_TARGET_LITTLEENDIAN is not ORed into flags then endianness the same as the host is assumed. If both DW_DLC_TARGET_LITTLEENDIAN and DW_DLC_TARGET_BIGENDIAN are OR-d in it is an error.

Either one of two output forms is specifiable:
DW_DLC_STREAM_RELOCATIONS or
DW_DLC_SYMBOLIC_RELOCATIONS .

The default is DW_DLC_STREAM_RELOCATIONS . The DW_DLC_STREAM_RELOCATIONS are relocations in a binary stream (as used in a MIPS/IRIX Elf object).

The DW_DLC_SYMBOLIC_RELOCATIONS are the same relocations but expressed in an array of structures defined by libdwarf, which the caller of the relevant function (see below) must deal with appropriately. This method of expressing relocations allows the producer-application to easily produce assembler text output of debugging information.

When DW_DLC_SYMBOLIC_RELOCATIONS is ORed into flags then relocations are returned not as streams but through an array of structures.

The function `func` must be provided by the user of this library. Its prototype is:

```
typedef int (*Dwarf_Callback_Func) (
    char* name,
    int size,
    Dwarf_Unsigned type,
    Dwarf_Unsigned flags,
    Dwarf_Unsigned link,
    Dwarf_Unsigned info,
    Dwarf_Unsigned* sect_name_index,
    void * user_data,
    int* error)
```

For each section in the object file that `libdwarf` needs to create, it calls this function once (calling it from `dwarf_transform_to_disk_form()`), passing in the section name, the section type, the section flags, the `link` field, and the `info` field. For an Elf object file these values should be appropriate Elf section header values. For example, for relocation callbacks, the `link` field is supposed to be set (by the app) to the index of the `syntab` section (the `link` field passed through the callback must be ignored by the app). And, for relocation callbacks, the `info` field is passed as the elf section number of the section the relocations apply to.

The `sect_name_index` field is a field you use to pass a symbol index back to `libdwarf`. In Elf, each section gets an elf symbol table entry so that relocations have an address to refer to (relocations rely on addresses in the Elf symbol table). You will create the Elf symbol table, so you have to tell `libdwarf` the index to put into relocation records for the section newly defined here.

On success the user function should return the Elf section number of the newly created Elf section.

On success, the function should also set the integer pointed to by `sect_name_index` to the Elf symbol number assigned in the Elf symbol table of the new Elf section. This symbol number is needed with relocations dependent on the relocation of this new section.

Use the `dwarf_producer_init_c()` interface instead of this interface.

For example, the `.debug_line` section's third data element (in a compilation unit) is the offset from the beginning of the `.debug_info` section of the compilation unit entry for this `.debug_line` set. The relocation entry in `.rel.debug_line` for this offset must have the relocation symbol index of the symbol `.debug_info` returned by the callback of that section-creation through the pointer `sect_name_index`.

On failure, the function should return -1 and set the `error` integer to an error code.

Nothing in `libdwarf` actually depends on the section index returned being a real Elf section. The Elf section is simply useful for generating relocation records. Similarly, the Elf symbol table index returned through the `sect_name_index` must be an index that can be used in relocations against this section. The application will probably want to note the values passed to this function in some form, even if no Elf file is being produced.

`errhand` is a pointer to a function that will be used as a default fall-back function for handling errors detected by `Libdwarf`.

`errarg` is the default error argument used by the function pointed to by `errhand`.

For historical reasons the error handling is complicated and the following three paragraphs describe the three possible scenarios when a producer function detects an error. In all cases a short error message is printed on `stdout` if the error number is negative (as all such should be, see `libdwarf.h`). Then further action is taken as follows.

First, if the `Dwarf_Error` argument to any specific producer function (see the functions documented below) is non-null the `errhand` argument here is ignored in that call and the specific producer function sets the `Dwarf_Error` and returns some specific value (for `dwarf_producer_init` it is `DW_DLV_OK` as mentioned just above) indicating there is an error.

Second, if the `Dwarf_Error` argument to any specific producer function (see the functions documented below) is `NULL` and the `errarg` to `dwarf_producer_init()` is non-`NULL` then on an error in the producer code the `Dwarf_Handler` function is called and if that called function returns the producer code returns a specific value (for `dwarf_producer_init` it is `DW_DLV_OK` as mentioned just above) indicating there is an error.

Third, if the `Dwarf_Error` argument to any specific producer function (see the functions documented below) is `NULL` and the `errarg` to `dwarf_producer_init()` is `NULL` then on an error `abort()` is called.

The `user_data` argument is not examined by `libdwarf`. It is passed to user code in all calls by `libdwarf` to the `Dwarf_Callback_Func()` function and may be used by consumer code for the consumer's own purposes. Typical uses might be to pass in a pointer to some user data structure or to pass an integer that somehow is useful to the `libdwarf`-using code.

The `isa_name` argument must be non-null and contain one of the strings defined in the `isa_relocs` array in `pro_init.c`: `"irix"`, `"mips"`, `"x86"`, `"x86_64"`, `"arm"`, `"arm64"`, `"ppc"`, `"ppc64"`, `"sparc"`. The names are not strictly ISA names (nor ABI names) but a hopefully-meaningful mixing of the concepts of ISA and ABI. The intent is mainly to define relocation codes applicable to `DW_DLC_STREAM_RELOCATIONS`. New `isa_name` values will be provided as users request. In the `"irix"` case a special relocation is defined so a special CIE reference field can be created (if and only if the augmentation string is `"z"`).

The `dwarf_version` argument should be one of `"V2"`, `"V3"`, `"V4"`, `"V5"` to indicate which DWARF version is the overall format to be emitted. Individual section version numbers will obey the standard for that overall DWARF version.

The `extra` argument supports a comma-separated list of options. Passing in a null pointer or an empty string is acceptable if no such options are needed or used. All-

lowercase option names are reserved to the libdwarf implementation itself (specific implementations may want to use a leading upper-case letter for additional options).

The available options are

```
"default_is_stmt",
"address_size",
"minimum_instruction_length",
"maximum_operations_per_instruction",
"opcode_base",
"line_base",
"line_range",
"linetable_version",
"segment_selector_size",
and
"segment_size".
```

For example, to set the line-table generation default value of `is_stmt` to 0 pass in

```
"default_is_stmt=0".
```

To also set the `minimum_instruction_length` used in calculating line table address-advance values to one one would pass in

```
"default_is_stmt=0,minimum_instruction_length=1".
```

It's appropriate to add

```
"opcode_base=13"
```

for DWARF3 through DWARF5. All these default to something, but the something depends on environment what macro names are set by the environment or a just constants which makes it difficult to alter these values. See `pro_line.h` for the use of line-table related constants (which will vary depending on the target ISA and ABI and compilers).

The `error` argument is set through the pointer to return specific error if `error` is non-null and there is an error. The error details will be passed back through this pointer argument.

5.1.2 `dwarf_pro_set_default_string_form()`

```
int dwarf_pro_set_default_string_form(
    Dwarf_P_Debug *dbg,
    int            desired_form,
    Dwarf_Error *error)
```

The function `dwarf_pro_set_default_string_form()` sets the `Dwarf_P_Debug` descriptor to favor one of the two allowed values: `DW_FORM_string` (the default) or `DW_FORM_strp`.

When `DW_FORM_strp` is selected very short names will still use form `DW_FORM_string`.

The function should be called immediately after a successful call to `dwarf_producer_init()`.

Strings for `DW_FORM_strp` are not duplicated in the `.debug_str` section: each unique string appears exactly once.

On success it returns `DW_DLV_OK`. On error it returns `DW_DLV_ERROR`.

5.1.3 `dwarf_transform_to_disk_form_a()`

```
int dwarf_transform_to_disk_form_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_Signed *chunk_count_out,  
    Dwarf_Error* error)
```

The function `dwarf_transform_to_disk_form_a()` is new in September 2016. It produces the same result as `dwarf_transform_to_disk_form()` but returns the count through the new pointer argument `chunk_count_out`.

On success it returns `DW_DLV_OK` and sets `chunk_count_out` to the number of chunks of section data to be accessed by `dwarf_get_section_bytes_a()`.

It turns the DIE and other information specified for this `Dwarf_P_Debug` into a stream of bytes for each section being produced. These byte streams can be retrieved from the `Dwarf_P_Debug` by calls to `dwarf_get_section_bytes_a()` (see below).

In case of error `dwarf_transform_to_disk_form_a()` returns `DW_DLV_ERROR`.

The number of chunks is used to access data by `dwarf_get_section_bytes_a()` (see below) and the section data provided your code will insert into an object file or the like. Each section of the resulting object is typically many small chunks. Each chunk has a section index and a length as well as a pointer to a block of data (see `dwarf_get_section_bytes_a()`).

For each unique section being produced `dwarf_transform_to_disk_form_a()` calls the `Dwarf_Callback_Func` exactly once. The callback provides the connection between Elf sections (which we presume is the object format to be emitted) and the `libdwarf()` internal section numbering.

For `DW_DLC_STREAM_RELOCATIONS` a call to `Dwarf_Callback_Func` is made by `libdwarf` for each relocation section. Calls to `dwarf_get_section_bytes_a()` (see below). allow the `dwarf_transform_to_disk_form_a()` caller to get byte streams and write them to an object file as desired, just as with the other sections of the object being created.

For `DW_DLC_SYMBOLIC_RELOCATIONS` the user code should use `dwarf_get_relocation_info_count()` and `dwarf_get_relocation_info()` to retrieve the relocation info generated by

`dwarf_transform_to_disk_form()` and do something with it.

On failure it returns `DW_DLV_ERROR` and returns an error pointer through `*error`.

5.1.3.1 `dwarf_transform_to_disk_form()`

```
Dwarf_Signed dwarf_transform_to_disk_form(  
    Dwarf_P_Debug dbg,  
    Dwarf_Error* error)
```

The function `dwarf_transform_to_disk_form()` is the original call to generate output and a better interface is used by `dwarf_transform_to_disk_form_a()` though both do the same work and have the same meaning.

5.1.4 `dwarf_get_section_bytes_a()`

```
int dwarf_get_section_bytes_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_Signed dwarf_section,  
    Dwarf_Signed *elf_section_index,  
    Dwarf_Unsigned *length,  
    Dwarf_Ptr *section_bytes,  
    Dwarf_Error* error)
```

The function `dwarf_get_section_bytes_a()` must be called repetitively, with the index `dwarf_section` starting at 0 and continuing for the number of sections returned by `dwarf_transform_to_disk_form_a()`.

It returns `DW_DLV_NO_ENTRY` to indicate that there are no more sections of Dwarf information. Normally one would index through using the `sectioncount` from `dwarf_transform_to_disk_form_a()` so `DW_DLV_NO_ENTRY` would never be seen.

For each successful return (return value `DW_DLV_OK`), `*section_bytes` points to `*length` bytes of data that are normally added to the output object in Elf section `*elf_section` by the producer application. It is illegal to call these in any order other than 0 through N-1 where N is the number of dwarf sections returned by `dwarf_transform_to_disk_form_a()`. The elf section number is returned through the pointer `elf_section_index`.

The `dwarf_section` number is ignored: the data is returned as if the caller passed in the correct `dwarf_section` numbers in the required sequence.

In case of an error, `DW_DLV_ERROR` is returned and the `error` argument is set to indicate the error.

There is no requirement that the section bytes actually be written to an elf file. For example, consider the `.debug_info` section and its relocation section (the call back function would result in assigning 'section' numbers and the link field to tie these together (`.rel.debug_info` would have a link to `.debug_info`). One could examine the relocations, split the `.debug_info` data at relocation boundaries, emit byte streams (in hex) as assembler output, and at each relocation point, emit an assembler directive with a symbol name for the assembler. Examining the relocations is awkward though. It is much better to use `dwarf_get_section_relocation_info()`

The memory space of the section byte stream is freed by the `dwarf_producer_finish_a()` call (or would be if the `dwarf_producer_finish_a()` was actually correct), along with all the other space in use with that `Dwarf_P_Debug`.

5.1.5 `dwarf_get_relocation_info_count()`

```
int dwarf_get_relocation_info_count(  
    Dwarf_P_Debug dbg,  
    Dwarf_Unsigned *count_of_relocation_sections ,  
    int *drd_buffer_version,  
    Dwarf_Error* error)
```

The function `dwarf_get_relocation_info()` returns, through the pointer `count_of_relocation_sections`, the number of times that `dwarf_get_relocation_info()` should be called.

The function `dwarf_get_relocation_info()` returns `DW_DLV_OK` if the call was successful (the `count_of_relocation_sections` is therefore meaningful, though `count_of_relocation_sections` could be zero).

`*drd_buffer_version` is the value 2. If the structure pointed to by the `*reldata_buffer` changes this number will change. The application should verify that the number is the version it understands (that it matches the value of `DWARF_DRD_BUFFER_VERSION` (from `libdwarf.h`)). The value 1 version was never used in production MIPS `libdwarf` (version 1 did exist in source).

It returns `DW_DLV_NO_ENTRY` if `count_of_relocation_sections` is not meaningful because `DW_DLC_SYMBOLIC_RELOCATIONS` was not passed to the `dwarf_producer_init_c()` `dwarf_producer_init_b()` or `dwarf_producer_init()` call (whichever one was used).

It returns `DW_DLV_ERROR` if there was an error, in which case `count_of_relocation_sections` is not meaningful.

5.1.6 dwarf_get_relocation_info()

```
int dwarf_get_relocation_info(  
    Dwarf_P_Debug dbg,  
    Dwarf_Signed *elf_section_index,  
    Dwarf_Signed *elf_section_index_link,  
    Dwarf_Unsigned *relocation_buffer_count,  
    Dwarf_Relocation_Data *reldata_buffer,  
    Dwarf_Error* error)
```

The function `dwarf_get_relocation_info()` should normally be called repetitively, for the number of relocation sections that `dwarf_get_relocation_info_count()` indicated exist.

It returns `DW_DLV_OK` to indicate that valid values are returned through the pointer arguments. The error argument is not set.

It returns `DW_DLV_NO_ENTRY` if there are no entries (the count of relocation arrays is zero.). The error argument is not set.

It returns `DW_DLV_ERROR` if there is an error. Calling `dwarf_get_relocation_info()` more than the number of times indicated by `dwarf_get_relocation_info_count()` (without an intervening call to `dwarf_reset_section_bytes()`) results in a return of `DW_DLV_ERROR` once past the valid count. The error argument is set to indicate the error.

Now consider the returned-through-pointer values for `DW_DLV_OK`.

`*elf_section_index` is the 'elf section index' of the section implied by this group of relocations.

`*elf_section_index_link` is the section index of the section that these relocations apply to.

`*relocation_buffer_count` is the number of array entries of relocation information in the array pointed to by `*reldata_buffer`.

`*reldata_buffer` points to an array of 'struct Dwarf_Relocation_Data_s' structures.

The version 2 array information is as follows:

```
enum Dwarf_Rel_Type {dwarf_drt_none,  
    dwarf_drt_data_reloc,  
    dwarf_drt_segment_rel,
```



```
        dwarf_drt_first_of_length_pair,  
        dwarf_drt_second_of_length_pair  
};  
typedef struct Dwarf_Relocation_Data_s * Dwarf_Relocation_Data;  
struct Dwarf_Relocation_Data_s {  
    unsigned char    drd_type; /* contains Dwarf_Rel_Type */  
    unsigned char    drd_length; /* typically 4 or 8 */  
    Dwarf_Unsigned    drd_offset; /* where the data to reloc is */  
    Dwarf_Unsigned    drd_symbol_index;  
};
```

The Dwarf_Rel_Type enum is encoded (via casts if necessary) into the single unsigned char drd_type field to control the space used for this information (keep the space to 1 byte).

The unsigned char drd_length field holds the size in bytes of the field to be relocated. So for elf32 object formats with 32 bit apps, drd_length will be 4. For objects with MIPS -64 contents, drd_length will be 8. For some dwarf 64 bit environments, such as ia64, drd_length is 4 for some relocations (file offsets, for example) and 8 for others (run time addresses, for example).

If drd_type is dwarf_drt_none, this is an unused slot and it should be ignored.

If drd_type is dwarf_drt_data_reloc this is an ordinary relocation. The relocation type means either (R_MIPS_64) or (R_MIPS_32) (or the like for the particular ABI. drd_length gives the length of the field to be relocated. drd_offset is an offset (of the value to be relocated) in the section this relocation stuff is linked to. drd_symbol_index is the symbol index (if elf symbol indices were provided) or the handle to arbitrary information (if that is what the caller passed in to the relocation-creating dwarf calls) of the symbol that the relocation is relative to.

When drd_type is dwarf_drt_first_of_length_pair the next data record will be drt_second_of_length_pair and the drd_offset of the two data records will match. The relevant 'offset' in the section this reloc applies to should contain a symbolic pair like

```
        .word    second_symbol - first_symbol  
to generate a length. drd_length gives the length of the field to be relocated.
```

drt_segment_rel means (R_MIPS_SCN_DISP) is the real relocation (R_MIPS_SCN_DISP applies to exception tables and this part may need further work). drd_length gives the length of the field to be relocated.

The memory space of the section byte stream is freed by the

`dwarf_producer_finish_a()` call (or would be if the `dwarf_producer_finish_a()` was actually correct), along with all the other space in use with that `Dwarf_P_Debug`.

5.1.7 `dwarf_reset_section_bytes()`

```
void dwarf_reset_section_bytes(  
    Dwarf_P_Debug dbg  
)
```

The function `dwarf_reset_section_bytes()` is used to reset the internal information so that `dwarf_get_section_bytes_a()` will begin (on the next call) at the initial dwarf section again. It also resets so that calls to `dwarf_get_relocation_info()` will begin again at the initial array of relocation information.

Some dwarf producers need to be able to run through the `dwarf_get_section_bytes_a()` and/or the `dwarf_get_relocation_info()` calls more than once and this call makes additional passes possible. The set of `Dwarf_Ptr` values returned is identical to the set returned by the first pass. It is acceptable to call this before finishing a pass of `dwarf_get_section_bytes_a()` or `dwarf_get_relocation_info()` calls. No errors are possible as this just resets some internal pointers. It is unwise to call this before `dwarf_transform_to_disk_form()` has been called.

5.1.8 `dwarf_pro_get_string_stats()`

```
int dwarf_pro_get_string_stats(  
    Dwarf_P_Debug dbg,  
    Dwarf_Unsigned * str_count,  
    Dwarf_Unsigned * str_total_length,  
    Dwarf_Unsigned * strp_count_debug_str,  
    Dwarf_Unsigned * strp_len_debug_str,  
    Dwarf_Unsigned * strp_reused_count,  
    Dwarf_Unsigned * strp_reused_len,  
    Dwarf_Error* error)
```

If it returns `DW_DLV_OK` the function `dwarf_pro_get_string_stats()` returns information about how `DW_AT_name` etc strings were stored in the output object. The values suggest how much string duplication was detected in the DWARF being created.

Call it after calling `dwarf_transform_to_disk_form()` and before calling `dwarf_producer_finish_a()`. It has no effect on the object being output.

On error it returns `DW_DLV_ERROR` and sets `error` through the pointer.

5.1.9 dwarf_producer_finish_a()

```
int dwarf_producer_finish_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_Error* error)
```

This is new in September 2016 and has the newer interface style, but is otherwise identical to `dwarf_producer_finish()` .

The function `dwarf_producer_finish_a()` should be called after all the bytes of data have been copied somewhere (normally the bytes are written to disk). It frees all dynamic space allocated for `dbg`, include space for the structure pointed to by `dbg`. This should not be called till the data have been copied or written to disk or are no longer of interest. It returns `DW_DLV_OK` if successful.

On error it returns `DW_DLV_ERROR` and sets `error` through the pointer.

5.2 Debugging Information Entry Creation

The functions in this section add new DIEs to the object, and also the relationships among the DIE to be specified by linking them up as parents, children, left or right siblings of each other. In addition, there is a function that marks the root of the graph thus created.

5.2.1 dwarf_add_die_to_debug_a()

```
int dwarf_add_die_to_debug_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die first_die,  
    Dwarf_Error *error)
```

The function `dwarf_add_die_to_debug_a()` indicates to Libdwarf the root DIE of the DIE graph that has been built so far. It is intended to mark the compilation-unit DIE for the object represented by `dbg`. The root DIE is specified by `first_die`.

It returns `DW_DLV_OK` on success, and `DW_DLV_error` on error.

5.2.2 dwarf_new_die_a()

```
int dwarf_new_die_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_Tag new_tag,  
    Dwarf_P_Die parent,  
    Dwarf_P_Die child,  
    Dwarf_P_Die left_sibling,  
    Dwarf_P_Die right_sibling,  
    Dwarf_P_Die *die_out,  
    Dwarf_Error *error)
```

On success `dwarf_new_die_a()` returns `DW_DLV_OK` and creates a new DIE with its parent, child, left sibling, and right sibling DIEs specified by `parent`, `child`, `left_sibling`, and `right_sibling`, respectively. The new die is passed to the caller via the argument `die_out()`. There is no requirement that all of these DIEs be specified, i.e. any of these descriptors may be `NULL`. If none is specified, this will be an isolated DIE. A DIE is transformed to disk form by `dwarf_transform_to_disk_form()` only if there is a path from the DIE specified by `dwarf_add_die_to_debug` to it.

The value of `new_tag` is the tag which is given to the new DIE. `parent`, `child`, `left_sibling`, and `right_sibling` are pointers to establish links to existing DIEs. Only one of `parent`, `child`, `left_sibling`, and `right_sibling` may be non-`NULL`. If `parent` (`child`) is given, the DIE is linked into the list after (before) the DIE pointed to. If `left_sibling` (`right_sibling`) is given, the DIE is linked into the list after (before) the DIE pointed to.

To add attributes to the new DIE, use the Attribute Creation functions defined in the next section.

On failure `dwarf_new_die_a()` returns `DW_DLV_ERROR` and sets `*error`.

5.2.3 dwarf_die_link_a()

```
int dwarf_die_link_a(  
    Dwarf_P_Die die,  
    Dwarf_P_Die parent,  
    Dwarf_P_Die child,  
    Dwarf_P_Die left_sibling,  
    Dwarf_P_Die right_sibling,  
    Dwarf_Error *error)
```

On success the function `dwarf_die_link_a()` returns `DW_DLV_OK` and links an existing DIE described by the given `die` to other existing DIEs. The given `die` can be linked to a parent DIE, a child DIE, a left sibling DIE, or a right sibling DIE by specifying non-`NULL` `parent`, `child`, `left_sibling`, and `right_sibling` `Dwarf_P_Die` descriptors.

Only one of `parent`, `child`, `left_sibling`, and `right_sibling` may be non-`NULL`. If `parent` (`child`) is given, the DIE is linked into the list after (before) the DIE pointed to. If `left_sibling` (`right_sibling`) is given, the DIE is linked into the list after (before) the DIE pointed to. Non-`NULL` links overwrite the corresponding links the given `die` may have had before the call to `dwarf_die_link_a()`.

If there is an error `dwarf_die_link_a()` returns `DW_DLV_ERROR` and sets `error` with the specific applicable error code.

5.3 DIE Markers

DIE markers provide a way for a producer to extract DIE offsets from DIE generation. The markers do not influence the generation of DWARF, they simply allow a producer to extract `.debug_info` offsets for whatever purpose the producer finds useful (for example, a producer might want some unique other section unknown to libdwarf to know a particular DIE offset).

One marks one or more DIEs as desired any time before calling `dwarf_transform_to_disk_form()`.

After calling `dwarf_transform_to_disk_form()` call `dwarf_get_die_markers()` which has the offsets where the marked DIEs were written in the generated `.debug_info` data.

5.3.1 dwarf_add_die_marker_a()

```
int dwarf_add_die_marker_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die die,  
    Dwarf_Unsigned marker,  
    Dwarf_Error *error)
```

This is preferred over `dwarf_add_die_marker()`. The function `dwarf_add_die_marker_a()` writes the value `marker` to the DIE descriptor given by `die`. Passing in a marker of 0 means 'there is no marker' (zero is the default in DIEs).

It returns `DW_DLV_OK`, on success. On error it returns `DW_DLV_ERROR`.

5.3.2 dwarf_get_die_marker_a()

```
int dwarf_get_die_marker_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die die,  
    Dwarf_Unsigned *marker,  
    Dwarf_Error *error)
```

The function `dwarf_get_die_marker_a()` returns the current marker value for this DIE through the pointer `marker`. A marker value of 0 means 'no marker was set'.

It returns `DW_DLV_OK`, on success. On error it returns `DW_DLV_ERROR`.

5.3.3 dwarf_get_die_markers_a()

```
int dwarf_get_die_markers_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Marker * marker_list,  
    Dwarf_Unsigned *marker_count,  
    Dwarf_Error *error)
```

The function `dwarf_get_die_markers_a()` returns a pointer to an array of `Dwarf_P_Marker` pointers to struct `Dwarf_P_Marker_s` structures through the pointer `marker_list`. The array length is returned through the pointer `marker_count`.

The call is only meaningful after a call to `dwarf_transform_to_disk_form()` as the transform call creates the struct `Dwarf_P_Marker_s` structures, one for each DIE generated for `.debug_info` (but only for DIEs that had a non-zero marker value). The field `ma_offset` in the structure is set during generation of the `.debug_info` byte stream. The field `ma_marker` in the structure is a copy of the DIE marker of the DIE given that offset.

It returns `DW_DLV_OK`, on success. On error it returns `DW_DLV_ERROR` (if there are no markers it returns `DW_DLV_ERROR`).

5.3.3.1 dwarf_get_die_markers()

```
Dwarf_Signed dwarf_get_die_markers(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Marker * marker_list,  
    Dwarf_Unsigned *marker_count,  
    Dwarf_Error *error)
```

The function `dwarf_get_die_marker()` returns a pointer to an array of `Dwarf_P_Marker` pointers to struct `Dwarf_P_Marker_s` structures through the pointer `marker_list`. The array length is returned through the pointer `marker_count`.

The call is only meaningful after a call to `dwarf_transform_to_disk_form()` as the transform call creates the struct `Dwarf_P_Marker_s` structures, one for each DIE generated for `.debug_info` (but only for DIEs that had a non-zero marker value). The field `ma_offset` in the structure is set during generation of the `.debug_info` byte stream. The field `ma_marker` in the structure is a copy of the DIE marker of the DIE given that offset.

It returns 0, on success. On error it returns `DW_DLV_BADADDR` (if there are no markers it returns `DW_DLV_BADADDR`).

5.4 Attribute Creation

The functions in this section add attributes to a DIE. These functions return a `Dwarf_P_Attribute` descriptor that represents the attribute added to the given DIE. In most cases the return value is only useful to determine if an error occurred.

Some of the attributes have values that are relocatable. They need a symbol with respect to which the linker will perform relocation. This symbol is specified by means of an index into the Elf symbol table for the object (of course, the symbol index can be more general than an index).

5.4.1 `dwarf_add_AT_location_expr_a()`

```
int dwarf_add_AT_location_expr_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attr,  
    Dwarf_P_Expr loc_expr,  
    Dwarf_P_Attr *attr_out,  
    Dwarf_Error *error)
```

On success the function `dwarf_add_AT_location_expr_a()` returns `DW_DLV_OK` and adds the attribute specified by `attr` to the DIE descriptor given by `ownerdie`. The new attribute is passed back to the caller through the pointer `attr_out`.

The attribute should be one that has a location expression as its value. The location expression that is the value is represented by the `Dwarf_P_Expr` descriptor `loc_expr`.

If the expression has a `DW_OP_addr` the code simply assumes that `DW_OP_addr` is the first operation and bases the only relocation that can be created on that assumption.

On error it returns `DW_DLV_ERROR`.

5.4.2 `dwarf_add_AT_name_a()`

```
int dwarf_add_AT_name_a(  
    Dwarf_P_Die ownerdie,  
    char *name,  
    Dwarf_P_Attribute * attr_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_AT_name_a()` adds the string specified by `name` as the value of the `DW_AT_name` attribute for the given DIE, `ownerdie`. It returns `DW_DLV_OK` on success and assigns the new attribute descriptor to `*attr_out`.

On error it returns DW_DLV_ERROR and does not set *attr_out.

5.4.3 dwarf_add_AT_comp_dir_a()

```
int dwarf_add_AT_comp_dir_a(  
    Dwarf_P_Die ownerdie,  
    char *current_working_directory,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function dwarf_add_AT_comp_dir_a adds the string given by current_working_directory as the value of the DW_AT_comp_dir attribute for the DIE described by the given ownerdie. On success it returns DW_DLV_OK and sets *attr_out to the new attribute.

On error, it returns DW_DLV_ERROR and does not touch attr_out .

5.4.4 dwarf_add_AT_producer_a()

```
int dwarf_add_AT_producer_a(  
    Dwarf_P_Die ownerdie,  
    char *producer_string,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function dwarf_add_AT_producer_a() adds the string given by producer_string as the value of the DW_AT_producer attribute for the DIE given by ownerdie.

On success it returns DW_DLV_OK and returns the new attribute descriptor representing this attribute through the pointer argument attr_out.

On error, it returns DW_DLV_ERROR.

5.4.5 dwarf_add_AT_any_value_sleb_a()

```
int dwarf_add_AT_any_value_sleb_a(  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attrnum,  
    Dwarf_Signed signed_value,  
    Dwarf_P_Attribute *out_attr,  
    Dwarf_Error *error)
```

The function dwarf_add_AT_any_value_sleb_a() adds the given Dwarf_Signed value signed_value as the value of the DW_AT_const_value attribute for the DIE described by the given ownerdie.

The FORM of the output value is DW_FORM_sdata (signed leb number) and the

attribute will be DW_AT_const_value.

On success it returns DW_DLV_OK and sets *out_attr to the created attribute.

On error, it returns DW_DLV_ERROR.

5.4.6 dwarf_add_AT_const_value_signedint_a()

```
int dwarf_add_AT_const_value_signedint_a(  
    Dwarf_P_Die ownerdie,  
    Dwarf_Signed signed_value,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function dwarf_add_AT_const_value_signedint_a adds the given Dwarf_Signed value signed_value as the value of the DW_AT_const_value attribute for the DIE described by the given ownerdie.

The FORM of the output value is DW_FORM_data<n> (signed leb number) and the attribute will be DW_AT_const_value.

With this interface and output, there is no way for consumers to know from the FORM that the value is signed.

On success it returns DW_DLV_OK and sets *attr_out to the created attribute.

On error, it returns DW_DLV_ERROR.

5.4.7 dwarf_add_AT_implicit_const()

```
int dwarf_add_AT_implicit_const(Dwarf_P_Die ownerdie,  
    Dwarf_Half attrnum,  
    Dwarf_Signed signed_value,  
    Dwarf_P_Attribute *outattr,  
    Dwarf_Error * error);
```

The function dwarf_add_AT_implicit_const creates a new attribute and adds the signed value to the abbreviation entry for this new attribute and attaches the new attribute to the DIE passed in.

The new attribute has

attrnum attribute for the DIE described by the given ownerdie. The form in the generated attribute is DW_FORM_implicit_const. The signed_value argument will be inserted in the abbreviation table as a signed leb value.

For a successful call the function returns DW_DLV_OK. and a pointer to the created argument is returned through the pointer outaddr.

In case of error the function returns DW_DLV_ERROR and no attribute is created.

5.4.8 dwarf_add_AT_any_value_uleb_a()

```
int dwarf_add_AT_any_value_uleb_a(  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attrnum,  
    Dwarf_Unsigned unsigned_value,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_AT_any_value_uleb_a` adds the given `Dwarf_Unsigned` value `unsigned_value` as the value of the `attrnum` attribute for the DIE described by the given `ownerdie`.

The FORM of the output value is DW_FORM_udata (unsigned leb number) and the attribute is `attrnum`.

On success it returns DW_DLV_OK and sets `*attr_out` to the newly created attribute.

On error, it returns DW_DLV_ERROR.

5.4.9 dwarf_add_AT_const_value_unsignedint_a()

```
int dwarf_add_AT_const_value_unsignedint_a(  
    Dwarf_P_Die ownerdie,  
    Dwarf_Unsigned unsigned_value,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_AT_const_value_unsignedint_a` adds the given `Dwarf_Unsigned` value `unsigned_value` as the value of the DW_AT_const_value attribute for the DIE described by the given `ownerdie`.

The FORM of the output value is DW_FORM_data<n> and the attribute will be DW_AT_const_value.

With this interface and output, there is no way for consumers to know from the FORM that the value is signed.

On success it returns DW_DLV_OK. and sets `*attr_out` to the newly created attribute.

On error, it returns DW_DLV_ERROR.

5.4.10 dwarf_add_AT_const_value_string_a()

```
int dwarf_add_AT_const_value_string_a(  
    Dwarf_P_Die ownerdie,  
    char *string_value,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_AT_const_value_string_a()` adds the string value given by `string_value` as the value of the `DW_AT_const_value` attribute for the DIE described by the given `ownerdie`.

On success it returns `DW_DLV_OK` *`attr_out` to a newly created attribute.

On error, it returns `DW_DLV_ERROR`.

5.4.11 dwarf_add_AT_targ_address_c()

```
int dwarf_add_AT_targ_address_c(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attr,  
    Dwarf_Unsigned pc_value,  
    Dwarf_Unsigned sym_index,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_AT_targ_address_c()` is identical to

`sym_index` is guaranteed to be large enough that it can contain a pointer to arbitrary data (so the caller can pass in a real elf symbol index, an arbitrary number, or a pointer to arbitrary data). The ability to pass in a pointer through `sym_index` is only usable with `DW_DLC_SYMBOLIC_RELOCATIONS`.

On success the function returns `DW_DLV_OK` `Dwarf_P_Attribute` and `pc_value` is put into the section stream output and the `sym_index` is applied to the relocation information.

On failure it returns `DW_DLV_ERROR`.

Do not use this function for attr `DW_AT_high_pc` if the value to be recorded is an offset (not a pc) [use `dwarf_add_AT_unsigned_const_afP` or `dwarf_add_AT_any_value_uleb_afP` instead].

On failure the function returns `DW_DLV_ERROR`

5.4.12 dwarf_add_AT_block_a()

```
int dwarf_add_AT_block_a(
    Dwarf_P_Debug      dbg,
    Dwarf_P_Die         ownerdie,
    Dwarf_Half          attr,
    Dwarf_Small         *block_data,
    Dwarf_Unsigned       block_size,
    Dwarf_P_Attribute*   attr_out,
    Dwarf_Error          *error)
```

This function works with all DW_FORM_block forms as well as DW_FORM_exprloc.

On success this returns DW_DLV_OK an attribute with a DW_FORM_block instance (does not create DW_FORM_block1, DW_FORM_block2, or DW_FORM_block4 at present) and returns a pointer to the new attribute through the pointer attr_out.

On failure this returns DW_DLV_ERROR

5.4.13 dwarf_add_AT_dataref_a()

```
int dwarf_add_AT_dataref_a(
    Dwarf_P_Debug dbg,
    Dwarf_P_Die ownerdie,
    Dwarf_Half attr,
    Dwarf_Unsigned pc_value,
    Dwarf_Unsigned sym_index,
    Dwarf_P_Attribute *attr_out,
    Dwarf_Error *error)
```

This is very similar to dwarf_add_AT_targ_address_b but results in a different FORM (results in DW_FORM_data4 or DW_FORM_data8).

Useful for adding relocatable addresses in location lists.

sym_index is guaranteed to be large enough that it can contain a pointer to arbitrary data (so the caller can pass in a real elf symbol index, an arbitrary number, or a pointer to arbitrary data). The ability to pass in a pointer through sym_index is only usable with DW_DLC_SYMBOLIC_RELOCATIONS.

On success it returns DW_DLV_OK and the pc_value is put into the section stream output and the sym_index is applied to the relocation information.

Do not use this function for DW_AT_high_pc, use dwarf_add_AT_unsigned_const or dwarf_add_AT_any_value_uleb [if the value to be recorded is an offset of DW_AT_low_pc] or dwarf_add_AT_targ_address_b [if the value to be recorded is an address].

5.4.14 dwarf_add_AT_ref_address_a

```
int dwarf_add_AT_ref_address_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attr,  
    Dwarf_Unsigned pc_value,  
    Dwarf_Unsigned sym_index,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

This is very similar to `dwarf_add_AT_targ_address_c` but results in a different FORM (results in `DW_FORM_ref_addr` being generated).

Useful for `DW_AT_type` and `DW_AT_import` attributes.

`sym_index()` is guaranteed to be large enough that it can contain a pointer to arbitrary data (so the caller can pass in a real elf symbol index, an arbitrary number, or a pointer to arbitrary data). The ability to pass in a pointer through `sym_index()` is only usable with `DW_DLC_SYMBOLIC_RELOCATIONS`.

On success the function returns `DW_DLV_OK` and `pc_value` is put into the section stream output and the `sym_index` is applied to the relocation information.

On failure the function returns `DW_DLV_ERROR`.

Do not use this function for `DW_AT_high_pc`.

5.4.15 dwarf_add_AT_unsigned_const_a()

```
int dwarf_add_AT_unsigned_const_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attr,  
    Dwarf_Unsigned value,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_AT_unsigned_const_a()` adds an attribute with a `Dwarf_Unsigned` value belonging to the "constant" class, to the DIE specified by `ownerdie`. The object that the DIE belongs to is specified by `dbg`. The attribute is specified by `attr`, and its value is specified by `value`.

The FORM of the output will be one of the `DW_FORM_data<n>` forms.

On success it returns DW_DLV_OK and sets *attr_out to the newly created attribute.

It returns DW_DLV_ERROR on error.

5.4.16 dwarf_add_AT_signed_const_a()

```
int dwarf_add_AT_signed_const_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attr,  
    Dwarf_Signed value,  
    Dwarf_P_Attribute *out_addr,  
    Dwarf_Error *error)
```

The function dwarf_add_AT_signed_const_a() adds an attribute with a Dwarf_Signed value belonging to the "constant" class, to the DIE specified by ownerdie. The object that the DIE belongs to is specified by dbg. The attribute is specified by attr, and its value is specified by value.

On success it returns DW_DLV_OK and sets *out_addr with a pointer to the new attribute.

On error it returns DW_DLV_ERROR.

5.4.17 dwarf_add_AT_reference_c()

```
int dwarf_add_AT_reference_c(  
    Dwarf_P_Debug dbg,  
    Dwarf_Half attr,  
    Dwarf_P_Die ownerdie,  
    Dwarf_P_Die otherdie,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

dwarf_add_AT_reference_c() accepts a NULL otherdie with the assumption that dwarf_fixup_AT_reference_die() will be called by user code to fill in the missing otherdie before the DIEs are transformed to disk form.

On success it returns DW_DLV_OK and returns a pointer to the new attribute through *attr_out.

On failure it returns DW_DLV_ERROR.

5.4.18 dwarf_fixup_AT_reference_die()

```
int dwarf_fixup_AT_reference_die(  
    Dwarf_Half attrnum,  
    Dwarf_P_Die ownerdie,  
    Dwarf_P_Die otherdie,  
    Dwarf_Error *error)
```

The function `dwarf_fixup_AT_reference_die()` is provided to set the NULL `otherdie` that `dwarf_add_AT_reference_c()` allows to the reference target DIE. This must be done before transforming to disk form. `attrnum()` should be the attribute number of the attribute of `Wownerdie` which is to be updated. For example, if a local forward reference was in a `WDW_AT_sibling` attribute in `ownerdie`, pass the value `WDW_AT_sibling` as `attrnum`.

Since no attribute number can appear more than once on a given DIE the `attrnum()` suffices to uniquely identify which attribute of `Wownerdie` to update

It returns either `DW_DLV_OK` (on success) or `DW_DLV_ERROR` (on error). Calling this on an attribute where `otherdie` was already set is an error.

5.4.19 dwarf_add_AT_flag_a()

```
int dwarf_add_AT_flag_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attr,  
    Dwarf_Small flag,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_AT_flag_a()` adds an attribute with a `Dwarf_Small` value belonging to the "flag" class, to the DIE specified by `ownerdie`. The object that the DIE belongs to is specified by `dbg`. The attribute is specified by `attr`, and its value is specified by `flag`.

On success it returns `DW_DLV_OK` and passes back a pointer to the new attribute through `*attr_out`.

On error it returns `DW_DLV_ERROR`.

5.4.20 dwarf_add_AT_string_a()

```
int dwarf_add_AT_string_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attr,  
    char *string,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_AT_string()` adds an attribute with a value that is a character string to the DIE specified by `ownerdie`. The object that the DIE belongs to is specified by `dbg`. The attribute is specified by `attr`, and its value is pointed to by `string`.

On success it returns `DW_DLV_OK` and set `*attr_out` with a pointer to the new attribute.

On failure it returns `DW_DLV_ERROR`.

5.4.21 `dwarf_add_AT_with_ref_sig8_a()`

```
int dwarf_add_AT_with_ref_sig8_a(  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attrnum,  
    const Dwarf_Sig8 *sig8_in,  
    Dwarf_P_Attribute *attr_out,  
    Dwarf_Error * error)
```

The function `dwarf_add_AT_with_sig8_a` creates an attribute containing the 8-byte signature block pointed to by `sig8_in` `DW_FORM_ref_sig8` with form `DW_FORM_ref_sig8`.

On success it returns `DW_DLV_OK` and sets `*attr_out` to the newly created attribute.

On failure it returns `DW_DLV_ERROR`.

5.4.22 `dwarf_add_AT_data16()`

```
int dwarf_add_AT_data16(  
    Dwarf_P_Die ownerdie,  
    Dwarf_Half attrnum,  
    Dwarf_Form_Data16 *ptr_to_val,  
    Dwarf_P_Attribute * attr_out,  
    Dwarf_Error * error)
```

The DWARF5 standard refers to 16 byte as simply data. It is up to the eventual reader of

the DWARF entry this call creates to understand what the sixteen bytes mean.

On success it returns DW_DLV_OK and returns the new attribute through the pointer `attr_out`.

On failure it returns DW_DLV_ERROR.

5.4.23 `dwarf_compress_integer_block()`

```
void* dwarf_compress_integer_block(  
    Dwarf_P_Debug    dbg,  
    Dwarf_Bool       unit_is_signed,  
    Dwarf_Small      unit_length_in_bits,  
    void*            input_block,  
    Dwarf_Unsigned    input_length_in_units,  
    Dwarf_Unsigned*   output_length_in_bytes_ptr,  
    Dwarf_Error*      error)
```

This was created in 2016 in support of the attribute DW_AT_SUN_func_offsets but the particular DWARF project involving this seems to have died. We have not provided a way to create the attribute. So this is pretty useless at this time.

5.5 Expression Creation

The following functions are used to convert location expressions into blocks so that attributes with values that are location expressions can store their values as a DW_FORM_blockn value. This is for both `.debug_info` and `.debug_loc` expression blocks.

To create an expression, first call `dwarf_new_expr_a()` to get a `Dwarf_P_Expr` descriptor that can be used to build up the block containing the location expression. Then insert the parts of the expression in prefix order (exactly the order they would be interpreted in in an expression interpreter). The bytes of the expression are then built-up as specified by the user.

5.5.1 `dwarf_new_expr_a()`

```
int dwarf_new_expr_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Expr *expr_out,  
    Dwarf_Error *error)
```

The function `dwarf_new_expr_a()` creates a new expression area in which a location expression stream can be created.

On success it returns DW_DLV_OK and returns a Dwarf_Expr Dwarf_Expr through the pointer which can be used to add operators a to build up a location expression.

On failure it returns DW_DLV_OK.

5.5.2 dwarf_add_expr_gen_a()

```
int dwarf_add_expr_gen_a(  
    Dwarf_P_Expr expr,  
    Dwarf_Small opcode,  
    Dwarf_Unsigned val1,  
    Dwarf_Unsigned val2,  
    Dwarf_Unsigned *stream_length_out,  
    Dwarf_Error *error)
```

The function dwarf_add_expr_gen_a() takes an operator specified by opcode, along with up to 2 operands specified by val1, and val2, converts it into the Dwarf representation and appends the bytes to the byte stream being assembled for the location expression represented by expr. The first operand, if present, to opcode is in val1, and the second operand, if present, is in val2. Both the operands may actually be signed or unsigned depending on opcode.

On success it returns DW_DLV_OK and sets *stream_length_out to the number of bytes in the byte stream for expr currently generated, i.e. after the addition of opcode.

It returns DW_DLV_ERROR on error.

The function dwarf_add_expr_gen_a() works for all opcodes except those that have a target address as an operand. This is because the function cannot not set up a relocation record that is needed when target addresses are involved.

5.5.3 dwarf_add_expr_addr_c()

```
int dwarf_add_expr_addr_c(  
    Dwarf_P_Expr expr,  
    Dwarf_Unsigned address,  
    Dwarf_Unsigned sym_index,  
    Dwarf_Unsigned *stream_length_out,  
    Dwarf_Error *error)
```

The function dwarf_add_expr_addr_c() is identical to dwarf_add_expr_addr_b() except that dwarf_add_expr_addr_c() returns a simple integer code.

sym_index() is guaranteed to be large enough that it can contain a pointer to arbitrary data (so the caller can pass in a real elf symbol index, an arbitrary number, or a pointer to arbitrary data). The ability to pass in a pointer through sym_index() is only

usable with DW_DLC_SYMBOLIC_RELOCATIONS.

On success the function returns DW_DLV_OK and sets *stream_length_out to the total length of the expression stream in expr.

On failure the function returns DW_DLV_ERROR.

5.5.4 dwarf_expr_current_offset_a()

```
int dwarf_expr_current_offset_a(  
    Dwarf_P_Expr expr,  
    Dwarf_Unsigned *stream_offset_out,  
    Dwarf_Error *error)
```

On success the function dwarf_expr_current_offset_a() returns DW_DLV_OK and sets *stream_offset_out to the current length in bytes of the expression stream.

On failure the function returns DW_DLV_ERROR.

5.5.5 dwarf_expr_into_block_a()

```
int dwarf_expr_into_block_a(  
    Dwarf_P_Expr expr,  
    Dwarf_Unsigned *length,  
    Dwarf_Small **address,  
    Dwarf_Error *error)
```

On success the function dwarf_expr_into_block_a() returns DW_DLV_OK and sets the length of the expr expression into *length and sets the value of a pointer into memory where the expression is currently held in the executing libdwarf into *address.

On failure it returns DW_DLV_ERROR.

5.5.6 dwarf_expr_reset()

```
void dwarf_expr_reset(  
    Dwarf_P_Expr expr,  
    Dwarf_Error *error)
```

This resets the expression content of expr() to be empty.

5.6 Line Number Operations

These are operations on the .debug_line section. They provide information about instructions in the program and the source lines the instruction come from. Typically,

code is generated in contiguous blocks, which may then be relocated as contiguous blocks. To make the provision of relocation information more efficient, the information is recorded in such a manner that only the address of the start of the block needs to be relocated. This is done by providing the address of the first instruction in a block using the function `dwarf_lne_set_address()`. Information about the instructions in the block are then added using the function `dwarf_add_line_entry_c()`, which specifies offsets from the address of the first instruction. The end of a contiguous block is indicated by calling the function `dwarf_lne_end_sequence()`.

Line number operations do not support `DW_DLC_SYMBOLIC_RELOCATIONS`.

5.6.1 `dwarf_add_line_entry_c()`

```
int dwarf_add_line_entry_c(  
    Dwarf_P_Debug dbg,  
    Dwarf_Unsigned file_index,  
    Dwarf_Addr code_offset,  
    Dwarf_Unsigned lineno,  
    Dwarf_Signed column_number,  
    Dwarf_Bool is_source_stmt_begin,  
    Dwarf_Bool is_basic_block_begin,  
    Dwarf_Bool is_epilogue_begin,  
    Dwarf_Bool is_prologue_end,  
    Dwarf_Unsigned isa,  
    Dwarf_Unsigned discriminator,  
    Dwarf_Error *error)
```

The function `dwarf_add_line_entry_c()` adds an entry to the section containing information about source lines. It specifies in `code_offset`, the address of this line. The function subtracts `code_offset` from the value given as the address of a previous line call to compute an offset, and the offset is what is recorded in the line instructions so no relocation will be needed on the line instruction generated.

The source file that gave rise to the instruction is specified by `file_index`, the source line number is specified by `lineno`, and the source column number is specified by `column_number` (column numbers begin at 1) (if the source column is unknown, specify 0). `file_index` is the index of the source file in a list of source files which is built up using the function `dwarf_add_file_decl()`.

`is_source_stmt_begin` is a boolean flag that is true only if the instruction at `code_address` is the first instruction in the sequence generated for the source line at `lineno`. Similarly, `is_basic_block_begin` is a boolean flag that is true only if the instruction at `code_address` is the first instruction of a basic block.

`is_epilogue_begin` is a boolean flag that is true only if the instruction at `code_address` is the first instruction in the sequence generated for the function

epilogue code.

Similarly, `is_prologue_end` is a boolean flag that is true only if the instruction at `code_address` is the last instruction of the sequence generated for the function prologue.

`isa` should be zero unless the code at `code_address` is generated in a non-standard isa. The values assigned to non-standard isas are defined by the compiler implementation.

`discriminator` should be zero unless the line table needs to distinguish among multiple blocks associated with the same source file, line, and column. The values assigned to `discriminator` are defined by the compiler implementation.

It returns `DW_DLV_OK` on success, and `DW_DLV_ERROR` on error.

5.6.2 `dwarf_lne_set_address_a()`

```
int dwarf_lne_set_address_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_Addr offs,  
    Dwarf_Unsigned symidx,  
    Dwarf_Error *error)
```

The function `dwarf_lne_set_address_a()` sets the target address at which a contiguous block of instructions begin. Information about the instructions in the block is added to `.debug_line` using calls to `dwarfdwarf_add_line_entry_c()` which specifies the offset of each instruction in the block relative to the start of the block. This is done so that a single relocation record can be used to obtain the final target address of every instruction in the block.

The relocatable address of the start of the block of instructions is specified by `offs`. The symbol used to relocate the address is given by `symidx`, which is normally the index of the symbol in the Elf symbol table.

It returns `DW_DLV_OK` on success, and `DW_DLV_ERROR` on error.

5.6.3 `dwarf_lne_end_sequence_a()`

```
int dwarf_lne_end_sequence_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_Addr address;  
    Dwarf_Error *error)
```

The function `dwarf_lne_end_sequence_a()` indicates the end of a contiguous block of instructions. `address()` should be just higher than the end of the last address in the sequence of instructions. Before the next block of instructions (if any) a call to

`dwarf_line_set_address_a()` will have to be made to set the address of the start of the target address of the block, followed by calls to `dwarf_add_line_entry_a()` for each of the instructions in the block.

It returns `DW_DLV_OK` on success and `DW_DLV_ERROR` on error.

5.6.4 `dwarf_add_directory_decl_a()`

```
int dwarf_add_directory_decl_a(  
    Dwarf_P_Debug dbg,  
    char *name,  
    Dwarf_Unsigned *index_in_directories,  
    Dwarf_Error *error)
```

The function `dwarf_add_directory_decl()` adds the string specified by name to the list of include directories in the statement program prologue of the `.debug_line` section. The string should therefore name a directory from which source files have been used to create the present object.

On success it returns `DW_DLV_OK` and sets the index of the string just added, in the list of include directories for the object. This index is then used to refer to this string. The index is passed back through the pointer argument `index_in_directories`

The first successful call of this function returns one, not zero, to be consistent with the directory indices that `dwarf_add_file_decl_a()` (below) expects.. DWARF5 is a bit different. TBD FIXME

It returns `DW_DLV_ERROR` on error.

5.6.5 `dwarf_add_file_decl_a()`

```
int dwarf_add_file_decl_a(  
    Dwarf_P_Debug dbg,  
    char *name,  
    Dwarf_Unsigned dir_idx,  
    Dwarf_Unsigned time_mod,  
    Dwarf_Unsigned length,  
    Dwarf_Unsigned *file_entry_count_out,  
    Dwarf_Error *error)
```

The function `dwarf_add_file_decl_a()` adds the name of a source file that contributed to the present object. The name of the file is specified by name (which must not be the empty string or a null pointer, it must point to a string with length greater than 0).

In case the name is not a fully-qualified pathname, it is considered prefixed with the name of the directory specified by `dir_idx` (which does not mean the name is changed or

physically prefixed by this producer function, we simply describe the meaning here). `dir_idx` is the index of the directory to be prefixed in the list buildup using `dwarf_add_directory_decl_a()`. As specified by the DWARF spec, a `dir_idx` of zero will be interpreted as meaning the directory of the compilation and another index must refer to a valid directory as FIXME

`time_mod` gives the time at which the file was last modified, and `length` gives the length of the file in bytes.

On success, it returns `DW_DLV_OK` and returns the index of the source file in the list built up so far through the pointer `file_entry_count_out`. This index can then be used to refer to this source file in calls to `dwarf_add_line_entry_a()`.

On error, it returns `DW_DLV_ERROR`.

5.7 Fast Access (aranges) Operations

These functions operate on the `.debug_aranges` section.

5.7.1 `dwarf_add_arange_c()`

```
int dwarf_add_arange_c(  
    Dwarf_P_Debug dbg,  
    Dwarf_Addr begin_address,  
    Dwarf_Unsigned length,  
    Dwarf_Unsigned symbol_index,  
    Dwarf_Unsigned end_symbol_index,  
    Dwarf_Addr offset_from_end_symbol,  
    Dwarf_Error *error)
```

The function `dwarf_add_arange_c()` adds another address range to be added to the section containing address range information, `.debug_aranges`.

If `end_symbol_index` is not zero we are using two symbols to create a length (must be `DW_DLC_SYMBOLIC_RELOCATIONS` to be useful)

`begin_address` is the offset from the symbol specified by `symbol_index`. `offset_from_end_symbol` is the offset from the symbol specified by `end_symbol_index`. `length` is ignored. This begin-end pair will be show up in the relocation array returned by `dwarf_get_relocation_info()` as a `dwarf_drt_first_of_length_pair` and `dwarf_drt_second_of_length_pair` pair of relocation records. The consuming application will turn that pair into something conceptually identical to

$$\text{.word end_symbol + offset_from_end - } \backslash \\ \text{(start_symbol + begin_address)}$$

The reason offsets are allowed on the begin and end symbols is to allow the caller to re-use existing labels when the labels are available and the corresponding offset is known (economizing on the number of labels in use). The 'offset_from_end - begin_address' will actually be in the binary stream, not the relocation record, so the app processing the relocation array must read that stream value into (for example) net_offset and actually emit something like

```
.word end_symbol - start_symbol + net_offset
```

If end_symbol_index is zero we must be given a length (either DW_DLC_STREAM_RELOCATIONS or DW_DLC_SYMBOLIC_RELOCATIONS):

The relocatable start address of the range is specified by begin_address, and the length of the address range is specified by length. The relocatable symbol to be used to relocate the start of the address range is specified by symbol_index, which is normally the index of the symbol in the Elf symbol table. The offset_from_end_symbol is ignored.

The function returns DW_DLV_OK on success and DW_DLV_ERROR on error.

5.8 DWARF5 .debug_sup section creation

The .debug_sup section (see the DWARF5 standard) enables symbolically linking two DWARF5 object files together.

5.8.1 dwarf_add_debug_sup()

This call provides all the information that the .debug_sup section has.

```
int dwarf_add_debug_sup(  
    Dwarf_P_Debug dbg,  
    Dwarf_Half     version,  
    Dwarf_Small     is_supplementary,  
    char           * filename,  
    Dwarf_Unsigned  checksum_len,  
    Dwarf_Small     * checksum,  
    Dwarf_Error *error)
```

On success it returns DW_DLV_OK and records the fields for creating the section.

The fields are as follows.

version should be passed in as 2.

filename must be a null-terminated string.

`is_supplementary` should be passed in as 0 or 1 depending on which type of object file is involved (see the DWARF5 standard).

`checksum` must be a byte array of length `checksum_len` used to validate (by a debugger) the use of the target object file.

`DW_DLV_NO_ENTRY` is never returned.

`DW_DLV_ERROR` is returned in case of an error, and `*error` is set as usual in `libdwarf`.

5.9 Fast Access (pubnames) Operations

These functions operate on the `.debug_pubnames` section.

5.9.1 `dwarf_add_pubname_a()`

```
int dwarf_add_pubname_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die die,  
    char *pubname_name,  
    Dwarf_Error *error)
```

It returns `DW_DLV_OK` on success and `DW_DLV_ERROR` on error.

5.10 Fast Access (pubtypes) Operations

These functions operate on the `.debug_pubtypes` section. An SGI-defined extension. Not part of standard DWARF.

5.10.1 `dwarf_add_pubtype_a()`

```
int dwarf_add_pubtype_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die die,  
    char *pubname_name,  
    Dwarf_Error *error)
```

It returns `DW_DLV_OK` on success and `DW_DLV_ERROR` on error.

5.11 Fast Access (weak names) Operations

These functions operate on the `.debug_weaknames` section. An SGI-defined extension. Not part of standard DWARF.

5.11.1 dwarf_add_weakname_a()

```
int dwarf_add_weakname_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die die,  
    char *weak_name,  
    Dwarf_Error *error)
```

It returns DW_DLV_OK on success and DW_DLV_ERROR on error.

5.12 Static Function Names Operations

The `.debug_funcnames` section contains the names of static function names defined in the object, and also the offsets of the DIES that represent the definitions of the functions in the `.debug_info` section. An SGI-defined extension. Not part of standard DWARF.

5.12.1 dwarf_add_funcname_a()

```
int dwarf_add_funcname_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die die,  
    char *func_name,  
    Dwarf_Error *error)
```

The function `dwarf_add_funcname_a()` adds the name of a static function specified by `func_name` to the section containing the names of static functions defined in the object represented by `dbg`. The DIE that represents the definition of the function is specified by `die`.

It returns DW_DLV_OK on success.

It returns DW_DLV_ERROR on error.

5.13 File-scope User-defined Type Names Operations

The `.debug_tynames` section contains the names of file-scope user-defined types in the given object, and also the offsets of the DIES that represent the definitions of the types in the `.debug_info` section. An SGI-defined extension. Not part of standard DWARF.

5.13.1 dwarf_add_typename_a()

```
int dwarf_add_typename_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die die,  
    char *type_name,  
    Dwarf_Error *error)
```

This the same as `dwarf_add_typename()` except that on success this returns DW_DLV_OK and on failure this returns DW_DLV_ERROR.

5.14 File-scope Static Variable Names Operations

The `.debug_varnames` section contains the names of file-scope static variables in the given object, and also the offsets of the DIEs that represent the definition of the variables in the `.debug_info` section. An SGI-defined section.

5.14.1 `dwarf_add_varname_a()`

```
int dwarf_add_varname_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Die die,  
    char *var_name,  
    Dwarf_Error *error)
```

This the same as `dwarf_add_varname()` except that on success this returns `DW_DLV_OK` and on failure this returns `DW_DLV_ERROR`.

5.15 Macro Information Creation

All strings passed in by the caller are copied by these functions, so the space in which the caller provides the strings may be ephemeral (on the stack, or immediately reused or whatever) without this causing any difficulty.

5.15.1 `dwarf_def_macro()`

```
int dwarf_def_macro(Dwarf_P_Debug dbg,  
    Dwarf_Unsigned lineno,  
    char *name  
    char *value,  
    Dwarf_Error *error);
```

Adds a macro definition. The name argument should include the parentheses and parameter names if this is a function-like macro. Neither string should contain extraneous whitespace. `dwarf_def_macro()` adds the mandated space after the name and before the value in the output DWARF section(but does not change the strings pointed to by the arguments). If this is a definition before any files are read, `lineno` should be 0. Returns `DW_DLV_ERROR` and sets `error` if there is an error. Returns `DW_DLV_OK` if the call was successful.

5.15.2 `dwarf_undef_macro()`

```
int dwarf_undef_macro(Dwarf_P_Debug dbg,  
    Dwarf_Unsigned lineno,  
    char *name,  
    Dwarf_Error *error);
```

Adds a macro un-definition note. If this is a definition before any files are read, `lineno` should be 0. Returns `DW_DLV_ERROR` and sets `error` if there is an error. Returns `DW_DLV_OK` if the call was successful.

5.15.3 `dwarf_start_macro_file()`

```
int dwarf_start_macro_file(Dwarf_P_Debug dbg,
                           Dwarf_Unsigned lineno,
                           Dwarf_Unsigned fileindex,
                           Dwarf_Error *error);
```

`fileindex` is an index in the `.debug_line` header: the index of the file name. See the function `dwarf_add_file_decl()`. The `lineno` should be 0 if this file is the file of the compilation unit source itself (which, of course, is not a `#include` in any file). Returns `DW_DLV_ERROR` and sets `error` if there is an error. Returns `DW_DLV_OK` if the call was successful.

5.15.4 `dwarf_end_macro_file()`

```
int dwarf_end_macro_file(Dwarf_P_Debug dbg,
                         Dwarf_Error *error);
```

Returns `DW_DLV_ERROR` and sets `error` if there is an error. Returns `DW_DLV_OK` if the call was successful.

5.15.5 `dwarf_vendor_ext()`

```
int dwarf_vendor_ext(Dwarf_P_Debug dbg,
                    Dwarf_Unsigned constant,
                    char *      string,
                    Dwarf_Error* error);
```

The meaning of the `constant` and the `string` in the macro info section are undefined by DWARF itself, but the string must be an ordinary null terminated string. This call is not an extension to DWARF. It simply enables storing macro information as specified in the DWARF document. Returns `DW_DLV_ERROR` and sets `error` if there is an error. Returns `DW_DLV_OK` if the call was successful.

5.16 Low Level (`.debug_frame`) operations

These functions operate on the `.debug_frame` section. Refer to `libdwarf.h` for the register names and register assignment mapping. Both of these are necessarily machine dependent.

5.16.1 dwarf_new_fde_a()

```
int dwarf_new_fde_a(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Fde *fde_out,  
    Dwarf_Error *error)
```

On success the function `dwarf_new_fde_a()` returns `DW_DLV_OK` and returns a pointer to the fde through `fde_out`. The descriptor should be used to build a complete FDE. Subsequent calls to routines that build up the FDE should use the same `Dwarf_P_Fde` descriptor.

It returns `DW_DLV_ERROR` on error.

5.16.2 dwarf_add_frame_cie_a()

```
int dwarf_add_frame_cie_a(  
    Dwarf_P_Debug dbg,  
    char *augmenter,  
    Dwarf_Small code_align,  
    Dwarf_Small data_align,  
    Dwarf_Small ret_addr_reg,  
    Dwarf_Ptr init_bytes,  
    Dwarf_Unsigned init_bytes_len,  
    Dwarf_Unsigned *cie_index_out,  
    Dwarf_Error *error);
```

On success The function `dwarf_add_frame_cie_a()` returns `DW_DLV_OK`, creates a CIE, and returns an index to it through the pointer `cie_index_out`.

CIEs are used by FDEs to setup initial values for frames. The augmentation string for the CIE is specified by `augmenter`. The code alignment factor, data alignment factor, and the return address register for the CIE are specified by `code_align`, `data_align`, and `ret_addr_reg` respectively. `init_bytes` points to the bytes that represent the instructions for the CIE being created, and `init_bytes_len` specifies the number of bytes of instructions.

There is no convenient way to generate the `init_bytes` stream. One just has to calculate it by hand or separately generate something with the correct sequence and use `dwarfdump -v` and `readelf` (or `objdump`) and some kind of hex dumper to see the bytes. This is a serious inconvenience!

On error it returns `DW_DLV_ERROR`.

5.16.3 dwarf_add_frame_fde_c()

```
int dwarf_add_frame_fde_c(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Fde fde,  
    Dwarf_P_Die die,  
    Dwarf_Unsigned cie,  
    Dwarf_Addr virt_addr,  
    Dwarf_Unsigned code_len,  
    Dwarf_Unsigned sym_idx,  
    Dwarf_Unsigned sym_idx_of_end,  
    Dwarf_Addr offset_from_end_sym,  
    Dwarf_Unsigned *index_to_fde,  
    Dwarf_Error* error)
```

This function is like `dwarf_add_frame_fde()` except that `dwarf_add_frame_fde_c()` has new arguments to allow use with `DW_DLC_SYMBOLIC_RELOCATIONS` and a new argument to return the fde index..

The function `dwarf_add_frame_fde_c()` adds the FDE specified by `fde` to the list of FDEs for the object represented by the given `dbg`.

`die` specifies the DIE that represents the function whose frame information is specified by the given `fde`. If the MIPS/IRIX specific `DW_AT_MIPS_fde` attribute is not needed in `.debug_info` pass in 0 as the `die` argument.

`cie` specifies the index of the CIE that should be used to setup the initial conditions for the given frame. `virt_addr` represents the relocatable address at which the code for the given function begins, and `sym_idx` gives the index of the relocatable symbol to be used to relocate this address (`virt_addr` that is). `code_len` specifies the size in bytes of the machine instructions for the given function.

If `sym_idx_of_end` is zero (may be `DW_DLC_STREAM_RELOCATIONS` or `DW_DLC_SYMBOLIC_RELOCATIONS`):

`virt_addr` represents the relocatable address at which the code for the given function begins, and `sym_idx` gives the index of the relocatable symbol to be used to relocate this address (`virt_addr` that is). `code_len` specifies the size in bytes of the machine instructions for the given function. `sym_idx_of_end` and `offset_from_end_sym` are unused.

If `sym_idx_of_end` is non-zero (must be `DW_DLC_SYMBOLIC_RELOCATIONS` to be useful):

`virt_addr` is the offset from the symbol specified by `sym_idx` .

`offset_from_end_sym` is the offset from the symbol specified by `sym_idx_of_end`. `code_len` is ignored. This begin-end pair will be show up in the relocation array returned by `dwarf_get_relocation_info()` as a `dwarf_drt_first_of_length_pair` and `dwarf_drt_second_of_length_pair` pair of relocation records. The consuming application will turn that pair into something conceptually identical to

```
.word end_symbol + begin - \  
      ( start_symbol + offset_from_end)
```

The reason offsets are allowed on the begin and end symbols is to allow the caller to re-use existing labels when the labels are available and the corresponding offset is known (economizing on the number of labels in use). The 'offset_from_end - begin_address' will actually be in the binary stream, not the relocation record, so the app processing the relocation array must read that stream value into (for example) `net_offset` and actually emit something like

```
.word end_symbol - start_symbol + net_offset
```

On success it returns `DW_DLV_OK` and returns index to the given `fde` through the pointer `index_to_fde`.

On error, it returns `DW_DLV_ERROR`.

5.16.4 `dwarf_add_frame_info_c()`

```
int dwarf_add_frame_info_c(  
    Dwarf_P_Debug    dbg,  
    Dwarf_P_Fde      fde,  
    Dwarf_P_Die       die,  
    Dwarf_Unsigned    cie,  
    Dwarf_Addr        virt_addr,  
    Dwarf_Unsigned    code_len,  
    Dwarf_Unsigned    sym_idx,  
    Dwarf_Unsigned    end_symbol_index,  
    Dwarf_Addr        offset_from_end_symbol,  
    Dwarf_Signed      offset_into_exception_tables,  
    Dwarf_Unsigned    exception_table_symbol,  
    Dwarf_Unsigned    *index_to_fde,  
    Dwarf_Error*      error)
```

On success The function `dwarf_add_frame_fde_c()` returns `DW_DLV_OK`, adds the FDE specified by `fde` to the list of FDEs for the object represented by the given `dbg`, and. passes the index of the `fde` back through the pointer `index_to_fde`

On failure it returns DW_DLV_ERROR.

5.16.5 dwarf_fde_cfa_offset_a()

```
int dwarf_fde_cfa_offset_a( Dwarf_P_Fde fde,  
    Dwarf_Unsigned reg,  
    Dwarf_Signed offset,  
    Dwarf_Error *error)
```

The function `dwarf_fde_cfa_offset()` appends a DW_CFA_offset operation to the FDE, specified by `fde`, being constructed. The first operand of the DW_CFA_offset operation is specified by `reg`. The register specified should not exceed 6 bits. The second operand of the DW_CFA_offset operation is specified by `offset`.

It returns DW_DLV_OK on success.

It returns DW_DLV_ERROR on error.

5.16.6 dwarf_add_fde_inst_a()

```
int dwarf_add_fde_inst_a(  
    Dwarf_P_Fde fde,  
    Dwarf_Small op,  
    Dwarf_Unsigned val1,  
    Dwarf_Unsigned val2,  
    Dwarf_Error *error)
```

The function `dwarf_add_fde_inst()` adds the operation specified by `op` to the FDE specified by `fde`. Up to two operands can be specified in `val1`, and `val2`. Based on the operand specified Libdwarf decides how many operands are meaningful for the operand. It also converts the operands to the appropriate datatypes (they are passed to `dwarf_add_fde_inst` as Dwarf_Unsigned).

It returns DW_DLV_OK on success.

It returns DW_DLV_ERROR on error.

5.16.7 dwarf_insert_fde_inst_bytes()

```
int dwarf_insert_fde_inst_bytes(  
    Dwarf_P_Debug dbg,  
    Dwarf_P_Fde fde,  
    Dwarf_Unsigned len,  
    Dwarf_Ptr ibytes,  
    Dwarf_Error *error)
```

The function `dwarf_insert_fde_inst_bytes()` inserts the byte array (pointed at by `ibytes` and of length `len`) of frame instructions into the fde `fde`. It is incompatible

with `dwarf_add_fde_inst()`, do not use both functions on any given `Dwarf_P_Debug`. At present it may only be called once on a given `fde`. The `len` bytes `ibytes` may be constructed in any way, but the assumption is they were copied from an object file such as is returned by the `libdwarf` consumer function `dwarf_get_fde_instr_bytes`.

It returns `DW_DLV_OK` on success, and `DW_DLV_ERROR` on error.

CONTENTS

1. INTRODUCTION	1
1.1 Copyright	1
1.2 Purpose and Scope	1
1.3 Document History	2
1.4 Definitions	2
1.5 Overview	3
1.6 Revision History	3
2. Type Definitions	4
2.1 General Description	4
2.2 Namespace issues	5
3. libdwarf and Elf and relocations	5
3.1 binary or assembler output	5
3.2 libdwarf relationship to Elf	5
3.3 libdwarf and relocations	6
3.4 symbols, addresses, and offsets	6
4. Memory Management	6
4.1 Read Only Properties	7
4.2 Storage Deallocation	7
4.3 Error Handling	7
5. Functional Interface	7
5.1 Initialization and Termination Operations	7
5.1.1 dwarf_producer_init()	8
5.1.2 dwarf_pro_set_default_string_form()	12
5.1.3 dwarf_transform_to_disk_form_a()	13
5.1.4 dwarf_get_section_bytes_a()	14
5.1.5 dwarf_get_relocation_info_count()	15
5.1.6 dwarf_get_relocation_info()	16
5.1.7 dwarf_reset_section_bytes()	18
5.1.8 dwarf_pro_get_string_stats()	18
5.1.9 dwarf_producer_finish_a()	19
5.2 Debugging Information Entry Creation	19
5.2.1 dwarf_add_die_to_debug_a()	19
5.2.2 dwarf_new_die_a()	19
5.2.3 dwarf_die_link_a()	20
5.3 DIE Markers	21

5.3.1	dwarf_add_die_marker_a()	21
5.3.2	dwarf_get_die_marker_a()	21
5.3.3	dwarf_get_die_markers_a()	21
5.4	Attribute Creation	23
5.4.1	dwarf_add_AT_location_expr_a()	23
5.4.2	dwarf_add_AT_name_a()	23
5.4.3	dwarf_add_AT_comp_dir_a()	24
5.4.4	dwarf_add_AT_producer_a()	24
5.4.5	dwarf_add_AT_any_value_sleb_a()	24
5.4.6	dwarf_add_AT_const_value_signedint_a()	25
5.4.7	dwarf_add_AT_implicit_const()	25
5.4.8	dwarf_add_AT_any_value_uleb_a()	26
5.4.9	dwarf_add_AT_const_value_unsignedint_a()	26
5.4.10	dwarf_add_AT_const_value_string_a()	27
5.4.11	dwarf_add_AT_targ_address_c()	27
5.4.12	dwarf_add_AT_block_a()	28
5.4.13	dwarf_add_AT_dataref_a()	28
5.4.14	dwarf_add_AT_ref_address_a	29
5.4.15	dwarf_add_AT_unsigned_const_a()	29
5.4.16	dwarf_add_AT_signed_const_a()	30
5.4.17	dwarf_add_AT_reference_c()	30
5.4.18	dwarf_fixup_AT_reference_die()	31
5.4.19	dwarf_add_AT_flag_a()	31
5.4.20	dwarf_add_AT_string_a()	31
5.4.21	dwarf_add_AT_with_ref_sig8_a()	32
5.4.22	dwarf_add_AT_data16()	32
5.4.23	dwarf_compress_integer_block()	33
5.5	Expression Creation	33
5.5.1	dwarf_new_expr_a()	33
5.5.2	dwarf_add_expr_gen_a()	34
5.5.3	dwarf_add_expr_addr_c()	34
5.5.4	dwarf_expr_current_offset_a()	35
5.5.5	dwarf_expr_into_block_a()	35
5.5.6	dwarf_expr_reset()	35
5.6	Line Number Operations	35
5.6.1	dwarf_add_line_entry_c()	36
5.6.2	dwarf_lne_set_address_a()	37
5.6.3	dwarf_lne_end_sequence_a()	37
5.6.4	dwarf_add_directory_decl_a()	38
5.6.5	dwarf_add_file_decl_a()	38

5.7	Fast Access (aranges) Operations	39
5.7.1	dwarf_add_arange_c()	39
5.8	DWARF5 .debug_sup section creation	40
5.8.1	dwarf_add_debug_sup()	40
5.9	Fast Access (pubnames) Operations	41
5.9.1	dwarf_add_pubname_a()	41
5.10	Fast Access (pubtypes) Operations	41
5.10.1	dwarf_add_pubtype_a()	41
5.11	Fast Access (weak names) Operations	41
5.11.1	dwarf_add_weakname_a()	42
5.12	Static Function Names Operations	42
5.12.1	dwarf_add_funcname_a()	42
5.13	File-scope User-defined Type Names Operations	42
5.13.1	dwarf_add_typename_a()	42
5.14	File-scope Static Variable Names Operations	43
5.14.1	dwarf_add_varname_a()	43
5.15	Macro Information Creation	43
5.15.1	dwarf_def_macro()	43
5.15.2	dwarf_undef_macro()	43
5.15.3	dwarf_start_macro_file()	44
5.15.4	dwarf_end_macro_file()	44
5.15.5	dwarf_vendor_ext()	44
5.16	Low Level (.debug_frame) operations	44
5.16.1	dwarf_new_fde_a()	45
5.16.2	dwarf_add_frame_cie_a()	45
5.16.3	dwarf_add_frame_fde_c()	45
5.16.4	dwarf_add_frame_info_c()	47
5.16.5	dwarf_fde_cfa_offset_a()	48
5.16.6	dwarf_add_fde_inst_a()	48
5.16.7	dwarf_insert_fde_inst_bytes()	48

A Producer Library Interface to DWARF

David Anderson

ABSTRACT

This document describes an interface to a library of functions to create DWARF debugging information entries and DWARF line number information. It does not make recommendations as to how the functions described in this document should be implemented nor does it suggest possible optimizations.

The document is oriented to creating DWARF version 2. Support for creating DWARF3 and DWARF4 and DWARF5 is only partial: various features since DWARF2 cannot be created.

Rev 4.2, 16 January 2021 0.3.4